

LARAVEL IN PRODUCTION

Architecture, Performance, and Operations at Scale



Nuruzzaman Milon

Laravel in Production

Architecture, Performance, and Operations at Scale

Nuruzzaman Milon

Copyright © 2026 Nuruzzaman Milon

All rights reserved.

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the author, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law.

This book was written in Markdown and produced with [Ibis Next](#), using mPDF for PDF generation and CommonMark for HTML rendering. Code samples use [0xProto](#) (SIL Open Font License 1.1).

First published: July 2026

To my family, for making every effort worthwhile.

Preface

Who this book is for

This book is written for a mid-to-senior Laravel engineer who already knows how to build a Laravel application and is now being asked to keep one alive - the person a growing team hands the architecture decisions, the performance incident, and the migration that has to happen without a maintenance window. It assumes you already know how routing, Eloquent, queues, and testing work in Laravel; it doesn't stop to explain the framework. Every page is instead about the layer of decisions that sits on top of it: how you structure a codebase as more than one team starts touching it, how you keep it fast once traffic outgrows what fits comfortably on one database, how you keep it correct when the same request can arrive twice, how you keep it available when a dependency you don't control goes down, and how you run all of that at 3 a.m. on a Tuesday without it becoming a story you're still telling a year later.

If you're looking for a Laravel tutorial - how to define a route, how to write a migration, how Eloquent relationships work - this isn't that book, and there are excellent ones already. If you're the engineer who's outgrown those books and is now staring at a monorepo, a queue backlog, or a balance that doesn't add up, keep reading.

How to read it

The book is organized into eight parts, and the parts are ordered by dependency rather than importance, so reading front to back works, but it isn't the only sane way through it. Part 1 (Foundations) and Part 2 (Performance Engineering) build vocabulary that later parts assume - "shared kernel," "N+1," "backfill," and so on. Part 5 (Observability & Monitoring) is worth reading earlier than its page number suggests, because "correlation ID," "SLI/SLO/SLA," and "burn rate" get reused heavily in Parts 4, 6, and 7. Beyond that, Parts 3 through 8 are largely independent of each other, so if you picked this book up because of a live incident with your queues, your database failover, or your CI pipeline, it's reasonable to skip straight to the relevant chapter. Chapter 34, the capstone, assumes you've read - or at least skimmed - everything before it: it's a single worked incident that touches nearly every decision in the book, and it's meant to be read last.

Each chapter follows the same shape. It opens with a concrete production failure - drawn from about thirteen years of real incidents, anonymized - then works through the pattern that would have prevented or fixed it, and closes with a short "Chapter recap" you can scan later without rereading the whole chapter. Not every one of those incidents happened in a Laravel project; over those years I've worked across many languages and frameworks, and where the original system wasn't Laravel, I've translated the failure into a Laravel setting so it can be demonstrated in the stack this book is about.

About the examples

Every incident described in this book actually happened in real production systems. The failure modes, the trade-offs, and the fixes are drawn from lived experience over thirteen years as a professional software engineer. What isn't real is any identifying detail: company names, product names, team names, exact numbers, and anything else that could point back to a specific employer or client has been changed, generalized, or recombined. The named fictional businesses that recur throughout the book - Northfield (an order-processing platform), Vaultly (a wallet service), Reelcast (a livestreaming and short-video app), and others - are stand-ins for this kind of company, not case studies of any single one; they exist because those are the systems that reliably produce the failure modes worth discussing, not because they map one-to-one onto any company you could name. Nothing here discloses a trade secret, a proprietary architecture, or confidential information belonging to any organization I've worked with; the anonymization exists specifically to protect that while keeping the underlying lesson intact. The architecture decisions and opinions in these pages are my own.

A note on unfamiliar terms

Abbreviations are spelled out the first time they're used in a chapter, but if you're dropping into a chapter out of order, Appendix A (Glossary) collects the recurring ones - SLO, SLA, SLI, MTTR, RPO/RTO, ADR, and the rest - in one place, each with a pointer back to the chapter that covers it properly.

```
-- What paginate()/simplePaginate() run for page 84,213 (offset
2,105,300):
SELECT * FROM orders ORDER BY created_at DESC, id DESC LIMIT 25 OFFSET
2105300;

-- What cursorPaginate() runs for the equivalent next page - a seek, not
a scan:
SELECT * FROM orders
WHERE (created_at, id) < ('2026-03-14 09:12:00', 84213)
ORDER BY created_at DESC, id DESC
LIMIT 25;
```

The offset version still has to walk (and discard) 2,105,300 rows before it can return the 25 that matter, even though the database never sends them to the application. The cursor version seeks directly to a known point via the `(created_at, id)` index and reads only the 25 rows it returns - which is why its cost stays flat whether the cursor points at row 25 or row 2,500,000.

The tradeoff is real: cursor pagination can't render "Page 3 of 84,213" or let a user jump straight to page 50, because there's no cheap way to know how many rows precede an arbitrary offset without counting them. For the admin panel, the fix was to drop the total-count display and page-number control in favor of next/previous cursor navigation - a small UX change that turned an eight-second page load back into a two-hundred-millisecond one.

Chapter recap

- Dev-scale seed data hides exactly the problems that matter at production scale - test query performance against realistically sized datasets, not two hundred seeded rows.
- Eager load relations (`with, loadMissing`) and project only needed columns (`select`) to eliminate both query-count and I/O blowups in loops;
`Model :: preventLazyLoading()` turns a missing eager load into a failing test instead of a slow endpoint three months later.
- Index the columns your queries actually filter and sort by, in the order they're used, and don't expect a low-selectivity column to earn its own index.
- `EXPLAIN (type, key, rows, Extra)` turns "this feels slow" into a specific, fixable finding - a full scan, a missing key, or an in-memory sort.

- Connection pooling (PgBouncer/ProxySQL) is what keeps a horizontally scaled fleet from exhausting the database's connection ceiling, independent of query volume.
- Offset-based pagination's `COUNT(*)` and deep-offset scans both degrade as a table grows; cursor pagination stays flat at the cost of losing total counts and arbitrary page jumps - a fair trade for large tables.

Chapter 8 - Safe Schema Evolution at Scale

The migration that only failed in production

On Northfield, a fictional order-processing platform, a routine cleanup task renames a column: `orders.buyer_email` becomes `orders.customer_email`, to match naming used everywhere else in the domain. The migration is small, obviously correct, and reviewed without objection. In staging, against a ten-thousand-row table, it runs in forty milliseconds. In production, against a table with two hundred million rows, the `ALTER TABLE` takes an exclusive lock for twenty-five minutes. Every write to `orders` blocks behind it. The API starts returning timeouts, the on-call engineer gets paged, and the migration that "obviously" worked has to be killed mid-run, leaving the schema in an ambiguous state.

Nothing about the migration was wrong, syntactically. It was tested at the wrong scale, the same failure mode as Chapter 7, applied to schema changes instead of queries. This chapter covers the patterns that make schema changes safe regardless of table size: non-blocking DDL, a backward-compatible rename pattern, batched backfills, and - the part teams forget most often - tracking who else reads your schema before you assume a "safe" change is actually safe for everyone.

Why "fast in staging" doesn't transfer

The cost of a schema change is a function of data volume, not query complexity, and that cost varies by database engine and by the specific operation:

- Adding a nullable column with no default is metadata-only and near-instant on both MySQL (InnoDB, modern versions) and PostgreSQL, regardless of table size.
- Adding a column with a default value, or adding a `NOT NULL` constraint to an existing column, historically required rewriting every row to populate or validate the new value - a cost proportional to table size. Recent versions of both engines have narrowed this (PostgreSQL avoids a full rewrite for constant defaults since version 11; MySQL's behavior depends on the specific `ALGORITHM` used), but the safe assumption is: verify the specific operation against the specific engine version in use, on a copy of production-sized data, before trusting it in staging.

- Renaming a column, dropping a column, or changing a column's type can each range from instant metadata changes to full table rewrites, again depending on engine and version.

The rule that survives all of that engine-specific detail: never trust a migration's timing from a small staging database. Run it against a production-sized copy (a recent snapshot, restored somewhere disposable) before it ships, and know in advance whether the specific operation takes a blocking lock.

Concretely, adding a **NOT NULL** constraint to `orders.customer_email` on PostgreSQL as a single step forces the database to scan and validate every one of two hundred million rows while holding a lock that blocks writes for the duration:

```
// Locks orders for the full validation scan - the whole table, all at
// once.
Schema::table('orders', function (Blueprint $table) {
    $table->string('customer_email')->nullable(false)->change();
});
```

Splitting it into "add the constraint without enforcing it yet" and "validate it separately" turns one long exclusive lock into a near-instant metadata change plus a scan that only takes a brief lock at the very end:

```
-- Step 1: instant - takes the constraint's existence for granted,
-- doesn't scan.
ALTER TABLE orders ADD CONSTRAINT customer_email_not_null
    CHECK (customer_email IS NOT NULL) NOT VALID;

-- Step 2: scans the table, but only takes a brief lock to flip the
-- constraint's state once the scan finds no violations - not for the
-- duration of the scan itself.
ALTER TABLE orders VALIDATE CONSTRAINT customer_email_not_null;
```

Run as two separate deploys, the first one is safe to ship immediately; the second can run whenever the backfill from the next section has finished, with no code depending on either step's timing.

Concurrent and non-blocking index creation

Adding an index to a large table is one of the most common migrations to get wrong, because a plain `CREATE INDEX` takes a lock that blocks writes for as long as the index build takes - which, on a two-hundred-million-row table, can be tens of minutes.

PostgreSQL's answer is `CREATE INDEX CONCURRENTLY`, which builds the index without holding a write lock, at the cost of a slower two-pass build and a real failure mode: if it's interrupted, it can leave behind an invalid index that has to be dropped and retried, not resumed.

```
CREATE INDEX CONCURRENTLY idx_orders_customer_created
ON orders (customer_id, created_at);
```

MySQL's InnoDB engine supports online DDL for many index operations via `ALGORITHM=INPLACE, LOCK=NONE`, which allows concurrent reads and writes during the build - but not every DDL operation supports every algorithm, and the safe habit is to check the specific operation against the running engine version rather than assume.

```
ALTER TABLE orders
ADD INDEX idx_orders_customer_created (customer_id, created_at),
ALGORITHM=INPLACE, LOCK=NONE;
```

Laravel's migration files don't manage this distinction automatically - a `Schema::table()` migration using `$table->index()` issues a plain `CREATE INDEX/ADD INDEX` unless the raw SQL or a database-specific driver option is used to request the non-blocking variant. On a large table, that's a decision to make explicitly, not a default to trust:

```

public function up(): void
{
    /*
     * Schema::table()→index() takes a blocking lock for the build -
    fine
     * on a small table, a production incident on a two-hundred-million-
    row one.
     */
    DB::statement(
        'CREATE INDEX CONCURRENTLY idx_orders_customer_created ON orders
(customer_id, created_at)'
    );
}

```

CREATE INDEX CONCURRENTLY can't run inside the transaction Laravel wraps each migration in by default, so the migration class also needs **public \$withinTransaction = false;** - forgetting that is the most common way this pattern fails silently in a migration file, even though it works fine when run by hand. And because an interrupted concurrent build leaves an unusable index behind rather than rolling back, always check for one before retrying:

```

-- Finds indexes left in an invalid state by an interrupted CONCURRENTLY
build.
SELECT indexrelid::regclass AS index_name
FROM pg_index
WHERE indisvalid = false;

-- Safe to drop and retry - an invalid index is never used by the
planner.
DROP INDEX CONCURRENTLY idx_orders_customer_created;

```

The backward-compatible column rename pattern

Renaming a column directly - **ALTER TABLE orders RENAME COLUMN buyer_email TO customer_email** - is fast as a pure metadata operation on most engines. The real danger

isn't the **ALTER** itself; it's that the moment it commits, every piece of code still referencing the old column name breaks atomically, everywhere, at once - including code you don't control (see the next section). The fix is to never do a rename as a single step. Expand it into phases that can each be shipped, verified, and rolled back independently:

Phase	Schema state	Application state	Rollback if something's wrong
1. Add	New column exists, nullable	Not yet used	Drop the new column - no impact
2. Dual-write	Both columns exist	App writes to both columns on every write path; reads still from the old column	Stop dual-writing - old column is still authoritative
3. Backfill	Both columns exist	Batch job populates the new column for existing rows	Backfill is idempotent - safe to re-run or pause
4. Migrate reads	Both columns exist	App reads from the new column; old column no longer read	Revert the read change - old column is still fully populated
5. Stop dual-write	Both columns exist	App writes only to the new column	Revert to dual-writing if any consumer still expects the old column
6. Drop	Old column removed	Fully migrated	Not reversible past this point - this phase only ships once nothing reads the old column

Each phase is a separate, independently deployable change. If phase 4 reveals a bug (the new column has a subtly different meaning than assumed), the rollback is "revert one deploy," not "recover from a completed rename." This is the same expand-and-contract shape used for the

shared-package upgrades in Chapter 6 - the pattern generalizes to any change where "old" and "new" need to coexist for a transition window.

Phase 1 is the only phase that touches the database schema directly - it's a plain, additive migration:

```
public function up(): void
{
    Schema::table('orders', function (Blueprint $table) {
        $table->string('customer_email')->nullable()->after('buyer_email');
        // BTW after() doesn't work in postgres
    });
}
```

Phase 2's dual-write lives in the model, not in a migration, and is the one line of application code every write path now depends on until phase 5 removes it:

```
protected static function booted(): void
{
    static::saving(function (Order $order) {
        if ($order->isDirty('buyer_email')) {
            $order->customer_email = $order->buyer_email;
        }
    });
}
```

Phase 6's drop is deliberately its own migration, shipped only after the checklist at the end of this chapter is fully checked off:

```
public function up(): void
{
    Schema::table('orders', function (Blueprint $table) {
        $table->dropColumn('buyer_email');
    });
}
```

Batching large backfills

Phase 3 above - populating the new column for two hundred million existing rows - is itself a hazard if done as one statement:

```
// Don't do this on a large table.
DB::table('orders')->update(['customer_email' =>
DB::raw('buyer_email')]);
```

A single **UPDATE** touching every row holds row locks (and, depending on engine and isolation level, can hold them for the duration of one very long transaction), competes with live traffic for the same rows, and - on replicated setups - can create significant replication lag, since a replica applies the same enormous write as one unit. The safe alternative is chunking: touch a bounded number of rows per statement, with a small pause between batches to let replicas catch up and to leave room for live traffic.

```

Order :: query()
  →whereNull('customer_email')
  →orderBy('id')
  →chunkById(1000, function ($orders) {
    foreach ($orders as $order) {
      $order→update(['customer_email' => $order→buyer_email]);
    }

    usleep(100_000); // 100ms - bound replication lag and lock
    contention
  });

```

Running this as a queued, resumable Artisan command (tracking the last processed ID) rather than a single migration step means it can be paused, monitored, and restarted without redoing completed work - important for a backfill that might run for hours against a genuinely large table:

```

final class BackfillCustomerEmail extends Command
{
    protected $signature = 'orders:backfill-customer-email {--fresh} {--dry-run}';

    public function handle(): int
    {
        $lastId = $this→option('fresh') ? 0 : (int)
Cache::get('backfill:customer_email:last_id', 0);
        $dryRun = $this→option('dry-run');
        $processed = 0;

        Order :: query()
            →whereNull('customer_email')
            →where('id', '>', $lastId)
            →orderBy('id')
            →chunkById(1000, function (Collection $orders) use
($dryRun, &$processed) {
                if (! $dryRun) {
                    foreach ($orders as $order) {
                        $order→update(['customer_email' =>

```

```

$order→buyer_email]);
    }

    /*
    * Persisted, not just held in memory - a restart
after a
    * deploy or an `artisan:down` picks up where this
left off
    * instead of re-scanning rows already backfilled.
    */
    Cache::forever('backfill:customer_email:last_id',
$order→last()→id);
    }

    $processed += $order→count();

    /*
    * Without this, a backfill that runs for hours looks
    * identical to a stuck one - there's nothing to
distinguish
    * "still working" from "silently hung" on the terminal.
    */
    $this→components→info(
        ($dryRun ? '[dry run] ' : '')."Backfilled
{$processed} rows so far (last ID: {$order→last()→id})."
    );

    usleep(100_000);
});

    $this→components→info($dryRun
        ? "Dry run complete: {$processed} rows would have been
backfilled."
        : 'Backfill complete: no rows remain with a null
customer_email.');
```

```

    return self::SUCCESS;
}
}

```

Dispatched via `Schedule :: command('orders:backfill-customer-email')->everyMinute()->withoutOverlapping()` (Chapter 11 covers scheduling and overlap guards in depth), this makes progress in small, bounded increments and survives being interrupted by a deploy at any point.

`--dry-run` is worth adding to every backfill command before it ever touches a large table: it walks the exact same query and chunking logic, reports how many rows and which ID range would be affected, and writes nothing - so a mistake in the `WHERE` clause shows up as a suspicious row count in dry-run output, not as two hundred million incorrectly updated rows. And it's always worth emitting a progress line per batch, not just a single message at the end - a backfill that runs for hours with no output is indistinguishable from one that's silently hung, and the per-batch line (row count so far, last ID reached) is what turns "is this still running?" into a question the terminal already answers.

Downstream consumers: your schema is a public API

The phased rename above protects the application. It does nothing for anything else that reads the same database directly and isn't part of the deploy you control: a nightly analytics export, another internal service with direct read access to the same tables (an anti-pattern, but a common inherited one), or a BI tool querying a read replica. From their perspective, dropping `buyer_email` in phase 6 is a breaking change shipped with zero notice, because nothing about the application's own migration process is aware they exist.

The fix starts with an inventory: before any schema change on a table of nontrivial age, identify who reads it besides the application - export jobs, other services, reporting queries, scheduled dashboards. For Northfield's rename, that inventory turned up an internal reporting export that ran nightly against `orders.buyer_email` directly, owned by a different team on a different release cycle.

Two patterns handle this without forcing every downstream consumer onto the application's timeline:

- **A compatibility view.** Keep a database view (or, during the dual-write window, the old column itself) that presents the old shape, so existing consumer queries keep working unmodified while the underlying table has already moved on. The view gets dropped only once every known consumer has confirmed it's no longer needed - not on the application's schedule, but on the slowest consumer's.

- **A deprecation window with an actual notification.** Whoever owns the schema change communicates the timeline to known downstream owners the same way a public API deprecates an endpoint: an announced window, not a silent drop. This only works if the inventory step actually happened - a consumer nobody knew about can't be notified.

Both patterns above are damage control for consumers that already read columns directly. The more durable fix, when the downstream consumer is an HTTP client rather than a direct database reader, is to never let that exposure exist in the first place: don't return `Model::toJson()` or a bare `$order->toArray()` from a controller, and don't let a resource's field names default to mirroring column names. A Laravel API Resource (or a `league/fractal` transformer, on a project that predates or doesn't use Resources) sits between the two and makes the wire format an explicit, independently versioned contract:

```
final class OrderResource extends JsonResource
{
    public function toArray(Request $request): array
    {
        return [
            'id' => $this->id,
            /*
             * The response field stays `buyer_email` - the contract
            every
             * existing API client already depends on - even though the
             * column behind it is `customer_email` after phase 6 ships.
             * The rename happens entirely on one side of this line.
            */
            'buyer_email' => $this->customer_email,
            'status' => $this->status,
            'total_cents' => $this->total_cents,
        ];
    }
}
```

With that boundary in place, the entire six-phase rename in this chapter can happen without a single HTTP API consumer noticing - the resource is the only place that has to change, and it changes in the opposite direction: mapping the new column to the response field the API already promised, rather than the API being what dictates the column name.

Bot and risk scoring: don't let a vendor outage become your outage

Many teams layer a third-party risk-scoring or bot-detection service (device fingerprinting, behavioral scoring, a managed CAPTCHA) on top of application-level rate limits. These are valuable, but they introduce a dependency that can fail independently of your own infrastructure - and a naive integration turns "the risk vendor is slow" into "nobody can log in."

The rule that matters here: **decide, explicitly, what happens when the risk check itself fails or times out**, rather than letting an exception bubble up and 500 the request. There are three defensible postures, and the right one depends on what the endpoint protects:

Failure mode	Fail closed (block)	Fail open (allow)
Behavior on vendor timeout	Treat as high-risk, block or challenge	Treat as unscored, let the request through
Right for	High-value, low-volume actions (large withdrawals, account recovery)	High-volume, low-marginal-risk actions (viewing a page, browsing a catalog)
Failure mode if wrong	Vendor blip becomes a full outage for real users	A vendor outage becomes a temporary abuse window
Mitigation	Short timeout (100-300ms) + circuit breaker so a slow vendor degrades to "unscored" quickly	Alert loudly on sustained fail-open so someone investigates rather than it becoming permanent

For Kindred's login endpoint, the pragmatic default is: apply the application-level rate limits (which are self-hosted and always available) unconditionally, and treat the risk score as an *additional* signal that can tighten the response (e.g. require a CAPTCHA) but should not be a hard dependency for the base throttle to function. That way a risk-vendor outage degrades the system to "slightly less sophisticated bot detection," not "login is down."

In code, that means a tight timeout on the vendor call and an explicit fail-open branch, rather than letting the exception propagate:

```
public function scoreLogin(string $ip, string $email): ?RiskScore
{
    try {
        $response = $this->http
            request
            →timeout(0.25) // 250ms: fail fast, don't hold up the login
            →post('https://risk-vendor.example.com/v1/score', [
                'ip' => $ip,
                'email' => $email,
            ]);

        return RiskScore::fromArray($response->json());
    } catch (ConnectionException|RequestException $e) {
        /*
         * Fail open: an unscored login still passes the application-
        level
         * throttle above, it only skips the extra CAPTCHA/challenge
        step.
         */
        Log::critical('Risk vendor unavailable, failing open',
            ['exception' => $e->getMessage()]);

        return null; // caller treats null as "unscored"
    }
}
```

The loud `Log::critical` call is what keeps a fail-open decision visible instead of silently permanent.

Backpressure: what a limit does after it triggers

A rate limit isn't just "reject the request." What you return, and what you do internally when a limit is being approached, is its own design surface:

- **Return 429 Too Many Requests**, not a generic 403 or 500 - well-behaved clients (and API consumers) know to back off on a 429 specifically, and a **Retry-After** header lets them do it without guessing.

Concretely, that's `RateLimiter::availableIn()` feeding straight into the response:

```
public function login(Request $request): JsonResponse
{
    $key = "login-ip:{$request->ip()}";

    if (RateLimiter::tooManyAttempts($key, 30)) {
        $retryAfter = RateLimiter::availableIn($key);

        return response()->json(
            ['message' => "Too many login attempts. Try again in
{$retryAfter} seconds."],
            429,
        )->header('Retry-After', $retryAfter);
    }

    // ...
}
```

- **Shed load progressively, not as a cliff.** If an upstream dependency (a payment gateway, a third-party API) is itself struggling, consider tightening your own limits proactively rather than waiting for it to start timing out request-by-request - this is the same backpressure idea covered for queues in Chapter 11, applied at the edge instead of the worker layer.
- **Distinguish abuse throttling from capacity throttling.** A login rate limit exists to stop abuse and should feel deliberately strict. A general API rate limit protecting your own database from being overwhelmed is a capacity control, and Chapter 15's load-testing work is what tells you where that number should actually sit - guessing it produces either a limit so loose it doesn't protect anything, or so tight it throttles normal peak traffic.

Chapter recap

- Rate limiting is a layered decision, not a single mechanism: pick an algorithm (fixed window for cheap coarse limits, token bucket for anything where burst behavior matters), a key (IP, user, action, or a composite), and a response (reject, challenge, or degrade).
- Composite keys (IP + username, device + account) catch attacks that defeat any single-axis limit, including distributed credential stuffing that no per-IP limit alone can stop.
- Increment attempt counters on failure, not on every request, or you'll rate-limit your legitimate power users.
- Third-party risk/bot signals are an enhancement, not a foundation - decide explicitly whether each check fails open or closed, and never let a vendor timeout become your outage.
- Return 429 with **Retry-After**, and treat abuse throttling (strict, security-motivated) and capacity throttling (informed by load testing, see Chapter 15) as two different tuning problems with two different failure costs.

Chapter 15 - Load Testing & Capacity Planning

The load test that passed, and the launch that didn't

Northfield, a fictional retail platform, is three weeks out from its biggest marketing event of the year - a flash sale expected to drive ten times normal traffic in a two-hour window. The team runs a load test: a script hammers the `/checkout` endpoint with a thousand concurrent requests, latency stays under 200ms, error rate stays near zero. Load test passed. Ship it.

On the day of the sale, the site falls over anyway - not on `/checkout`, but on the *combination* of product-listing traffic, cart updates, and checkout happening together, because real shoppers browse, add to cart, abandon, come back, and only a fraction ever reach checkout at all. The single-endpoint test never exercised that mix, never exercised the database connection pool under concurrent reads-and-writes, and never came close to discovering that the actual constraint that day is the third-party payment gateway's own rate limit, which throttles Northfield's checkout calls once volume passes a threshold nobody had checked.

Nothing about that load test was fake or dishonest. It was just testing the wrong thing: one endpoint, one traffic shape, with no predefined target for what "passing" even meant beyond "didn't 500." This chapter is about closing both gaps - designing a load test that resembles the traffic you'll actually get, and deciding in advance, in the calm of a planning meeting, what numbers count as success.

Define the SLOs before you touch a load-testing tool

A Service Level Objective (SLO) is a target for a Service Level Indicator (SLI) - a specific, measurable signal like "95th percentile response time" or "checkout error rate." Without SLOs written down first, a load test has no way to fail: any result can be read as "fine," because there was never a bar to clear.

Write these down before running anything:

- **Latency targets, as percentiles, not averages.** "p50 < 150ms, p95 < 500ms, p99 <

1500ms" tells you far more than "average response time is 200ms," because an average hides the tail - and the tail is what users and monitoring alerts actually notice. A system with a 200ms average and a 8-second p99 is failing for one in a hundred users on every single request.

- **Error rate ceiling.** "Error rate stays under 0.5% at target load" - and be explicit about what counts as an error: 5xx responses, certainly, but also timeouts, and arguably 429s if the plan didn't intend to be rate-limiting real users at that volume.
- **Throughput target, grounded in a real number.** Not "handle a lot of traffic," but "sustain 2,000 checkout completions per minute for a 90-minute window," derived from the actual marketing forecast (expected visitors x historical conversion rate), not a round number that sounded safe in a meeting.

These three - latency, error rate, throughput - are the core of almost every load test's pass/fail criteria, and writing them down first forces the useful question: *what will we do if we hit the target load and any one of these fails?* If nobody can answer that in the planning meeting, the load test is going to surface a decision nobody's ready to make, under worse time pressure, on the day it matters.

Traffic shape: model the journey, not the endpoint

Real traffic is not N identical requests to one URL. A realistic load test reproduces the *mix* and *sequencing* of what users actually do, because bottlenecks frequently live in the interaction between paths, not any single path in isolation.

For the flash-sale scenario, a defensible traffic model looks like a weighted set of user journeys rather than a single script:

```

{
  "journeys": [
    { "name": "browse_only", "weight": 55, "steps": ["home", "category",
"product"] },
    { "name": "browse_and_abandon_cart", "weight": 25, "steps":
["product", "add_to_cart", "cart"] },
    { "name": "full_checkout", "weight": 15, "steps": ["product",
"add_to_cart", "cart", "checkout", "payment"] },
    { "name": "returning_search", "weight": 5, "steps": ["search",
"product", "add_to_cart", "checkout", "payment"] }
  ],
  "think_time_ms": { "min": 500, "max": 4000 },
  "ramp": { "from_rps": 50, "to_rps": 2000, "over_minutes": 15 }
}

```

Two details matter as much as the journey mix itself. First, **think time** - the pause between a real user's clicks - matters because it changes how many requests are in flight concurrently versus how many total requests land over a window; a script with zero think time can generate the right total volume while wildly overstating concurrent connection pressure (or understating it, depending on what you're testing for). Second, the **ramp**, not just the peak - a flash sale doesn't teleport to 2,000 requests per second, and a system that survives a 15-minute ramp to peak can still fail against a 90-second spike to the same peak, because caches haven't warmed and autoscaling hasn't caught up. If spiky arrival is plausible for your event, test that shape specifically (see the load-test type table below) rather than assuming a gentle ramp covers it.

The **full_checkout** journey above turns directly into a k6 scenario - think time between each step, and a **ramping-vus** (virtual users - k6's term for simulated concurrent users, each running the script below in its own loop) stage that mirrors the model's ramp. k6 (from Grafana Labs) is the tool used throughout this chapter's examples: a load-testing tool where the test script itself is JavaScript, run from the command line via **k6 run checkout.js**:

```

// k6 (https://k6.io) - the load-testing tool this chapter's examples
use.
import http from 'k6/http';
import { sleep } from 'k6';

export const options = {
  scenarios: {
    full_checkout: {
      executor: 'ramping-vus',
      startVUs: 0,
      stages: [
        { duration: '15m', target: 300 }, // ramp to ~15% of the 2,000
rps peak
        { duration: '90m', target: 300 }, // hold for the sale window
      ],
    },
  },
};

export default function () {
  http.get('https://shop.example.com/product/1234');
  sleep(Math.random() * 3.5 + 0.5); // think time, matching
think_time_ms above
  http.post('https://shop.example.com/cart/add', { product_id: 1234 });
  sleep(Math.random() * 3.5 + 0.5);
  http.get('https://shop.example.com/checkout');
  sleep(Math.random() * 3.5 + 0.5);
  http.post('https://shop.example.com/payment', { method: 'card' });
}

```

Each of the other journeys in the model would get its own scenario, weighted to its share of the mix.

Four kinds of load test, and what each one is actually for

This is a sample from "Laravel in Production: Architecture, Performance, and Operations at Scale".